



Linux For Embedded Systems

For Arabs

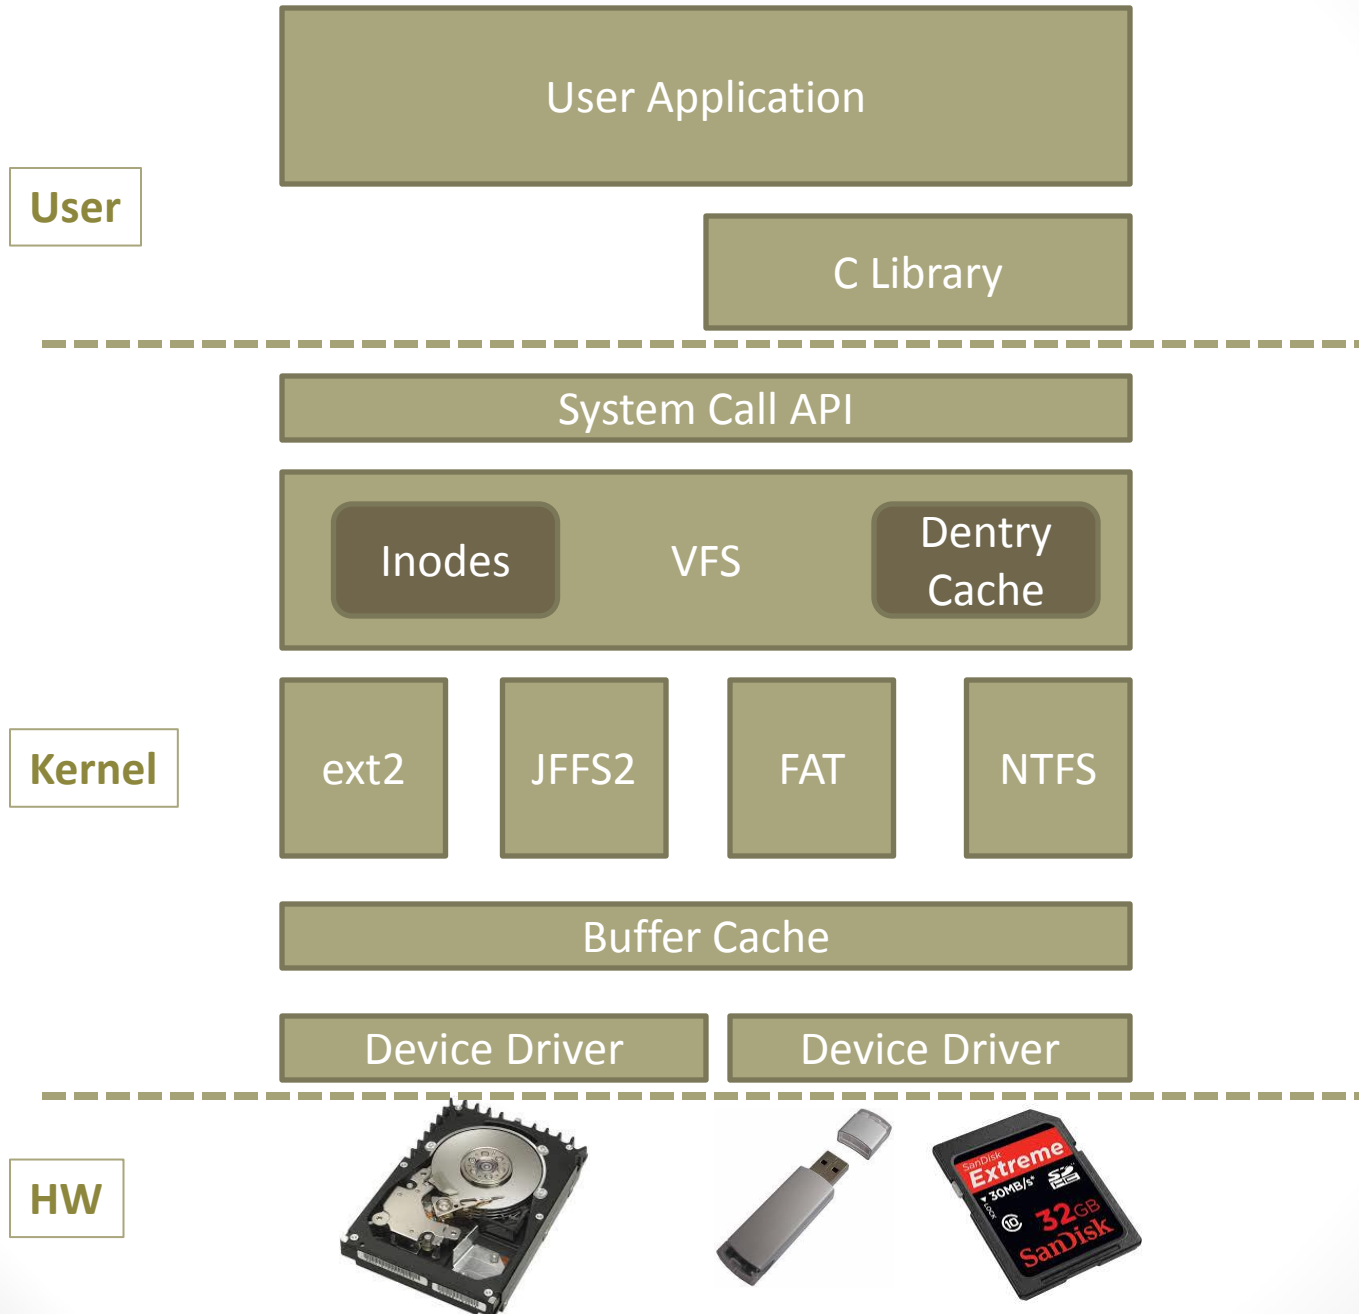
Course 102: Understanding Linux

Ahmed ElArabawy



Lecture 27:

FileSystems in Linux (Part 2)



The Big Picture

- User applications access files using a standard API
- This does not change no matter what storage media, partition, or used file system type
- The used API results in a system call towards the kernel (generated by the C-Library)
- The system call is received by the VFS (Virtual File System) subsystem of the Linux Kernel
- This subsystem provides a unified access to the file via the concepts of nodes and dentries (A previous lecture covered VFS in details)
- VFS then communicate with the used FileSystem that contain this file
- The filesystem in turn Communicate with the device driver for the storage media using the block abstraction
- The device driver is responsible for interfacing with the storage hardware device

Popular FileSystem Types

Linux Extended FileSystems



ext2

- The **ext2** used to be the most popular filesystem for Linux distributions running on a hard disk
- It uses block sizes of 1024, 2048, or 4096 Bytes
 - Large values waste disk space if we have a lot of small files
 - Small values increase overhead for metadata to point to the different blocks in the case of large files
- A problem with the ext2 was the risk of filesystem corruption due to unexpected reboots or power failures
 - A planned reboot will unmount all filesystems before the system is shutdown
 - This guarantees that the shutdown does not happen in the middle of a write operation
 - However, in case of a sudden power failure or a system crash, the filesystem will not be unmounted, and we may run into a corruption of the data in the fileSystem (some of the data is written without updating the filesystem metadata)
- The solution for this problem is “**Journaling**”

Journaling

- The filesystem that support journaling will have a special file ***“The Journal File”***
- Every time there is a modification to the filesystem (a write operation), this change is tracked first in the journal file then committed to the actual file in the filesystem
- This way, we will have a journaling points that we can revert to in case of a corruption (mismatch between filesystem contents and its meta-data)
- Upon a boot after a sudden shutdown/reboot, the journal file is tracked, and compared to the contents of the filesystem
- Changes in the journal may be applied or removed to maintain data consistency
- Most newer filesystems support Journaling to maintain the filesystem contents data integrity

Linux Extended FileSystems



ext3

- The **ext3** filesystem is an extension to the **ext2** filesystem to support Journalling
- It is both forward and backward compatible with **ext2** (we can convert **ext2** filesystem to **ext3** and vice versa)
- In **ext2** when the system shuts down abruptly, the next boot takes long time, because a consistency check is run on the filesystem
- In **ext3**, no consistency check is needed, the journaling file is checked to verify changes, which is a much faster process

Linux Extended FileSystems



ext4

- The **ext4** filesystem is an extension to the **ext3** filesystem
- Currently, it is the default filesystem for Linux
- It also supports journaling
- It removes some of the limitations of **ext3**,
 - The **ext4** filesystem can support filesystem size of more than the 16 terabyte which is the limit for **ext3**
 - The **ext4** filesystem can support file sizes of up to 1 Terabyte

Second Generation Journaling FileSystem

JFFS2 FileSystems



- The jffs2 filesystem is used with flash memory storage devices
- Use of flash memory storage is very common in embedded systems
- Flash memory has the following features:
 - Multiple files per Block:
 - The block size is large (tens to hundreds of kilobytes). A typical value is 128KB
 - This means, we need to be able to store multiple files in the same block
 - Accordingly, one block on the flash may contain several small files
 - Slow write operations
 - Writing a value to an empty place in the flash is performed one byte (or word) at a time (as normal devices)
 - However, you can not erase (or modify) a single byte (word)
 - Erasing requires the whole block to be erased
 - This means, if we need to modify a small file, this will require the whole block to be erased, then the block (or another empty one) will be re-written
 - This makes write operations in the flash much slower than other devices
 - This results in a higher chance of corruption due to power failures during a write operation
 - Flash Lifetime
 - Another limitation, flash memory has limited lifetime (specified in number of write operations)
 - Flash memory life time is measured by the number of write operations
 - We need to even out the write operations to avoid damaging parts of the flash too early

Second Generation Journaling FileSystem

JFFS2 FileSystems



- The jffs2 filesystem handles the flash as follows,
 - Due to the slow write operation, journaling becomes very essential, and hence journaling is supported in JFFS2
 - It makes sure that blocks are used evenly to distribute the write operations on the flash. This is called **Wear Leveling**
- Special care is needed with the use of tools that keep updating files such as logging tools (ex. *syslogd*, and *klogd*)
 - If we have other forms of storage, it would be better to direct the output away from the flash memory
 - Use of caching may be useful to reduce the number of modifications
 - This affects both system performance and flash life-time

Cram FileSystem

cramfs



- This filesystem objective is to compress a filesystem into a small ROM
- It is a read-only filesystem
- It supports compression of the data in the filesystem
- Useful for small embedded systems to store read only data in a small ROM or flash
- Ideal for boot ROMs

Memory Hosted FileSystems



ramfs

- A filesystem that lives in the system memory (RAM)
- This provides high access speed, but it is volatile (erased at reboot time or at shutdown)
- It is different from RamDisks that it can grow and shrink based on the need


Memory Hosted FileSystems



tmpfs

- Like **ramfs**, everything is stored in system volatile memory (RAM)
- Accordingly, contents of this filesystem are lost on power failure or reboot
- Different from the ramfs in that it can not grow dynamically
- It can also use the swap while the ramfs can not
- Normally mount to ***/tmp***

Network File System



Network FileSystem

NFS



- This filesystem will exist on a remote machine and will be accessed through the network
- Useful for sharing folders in the network
 - A central NFS server will contain the filesystem data
 - All machines that need to have access to the data will need to contain a NFS Client
 - Machines with NFS Clients will need to mount the NFS filesystem (Map to the network drive)

NFS in Embedded Systems



- Another very useful application for NFS is development of Embedded Systems
 - Embedded target flash memory size may not be able to hold all the tools and utilities used during development
 - Hence, all the tools and utilities can be located on a remote machine (NFS Server) and the target would mount an NFS filesystem to have access to it
 - Also, during development, we don't need to upload the image of the binary everytime we make a new build, instead, we keep the binary on the development host machine
 - The target can even have its root file system mounted as NFS, so the target will only carry the bootloader and kernel. All the rest will be on the remote machine (development host)

Virtual FileSystems

procfs & sysfs



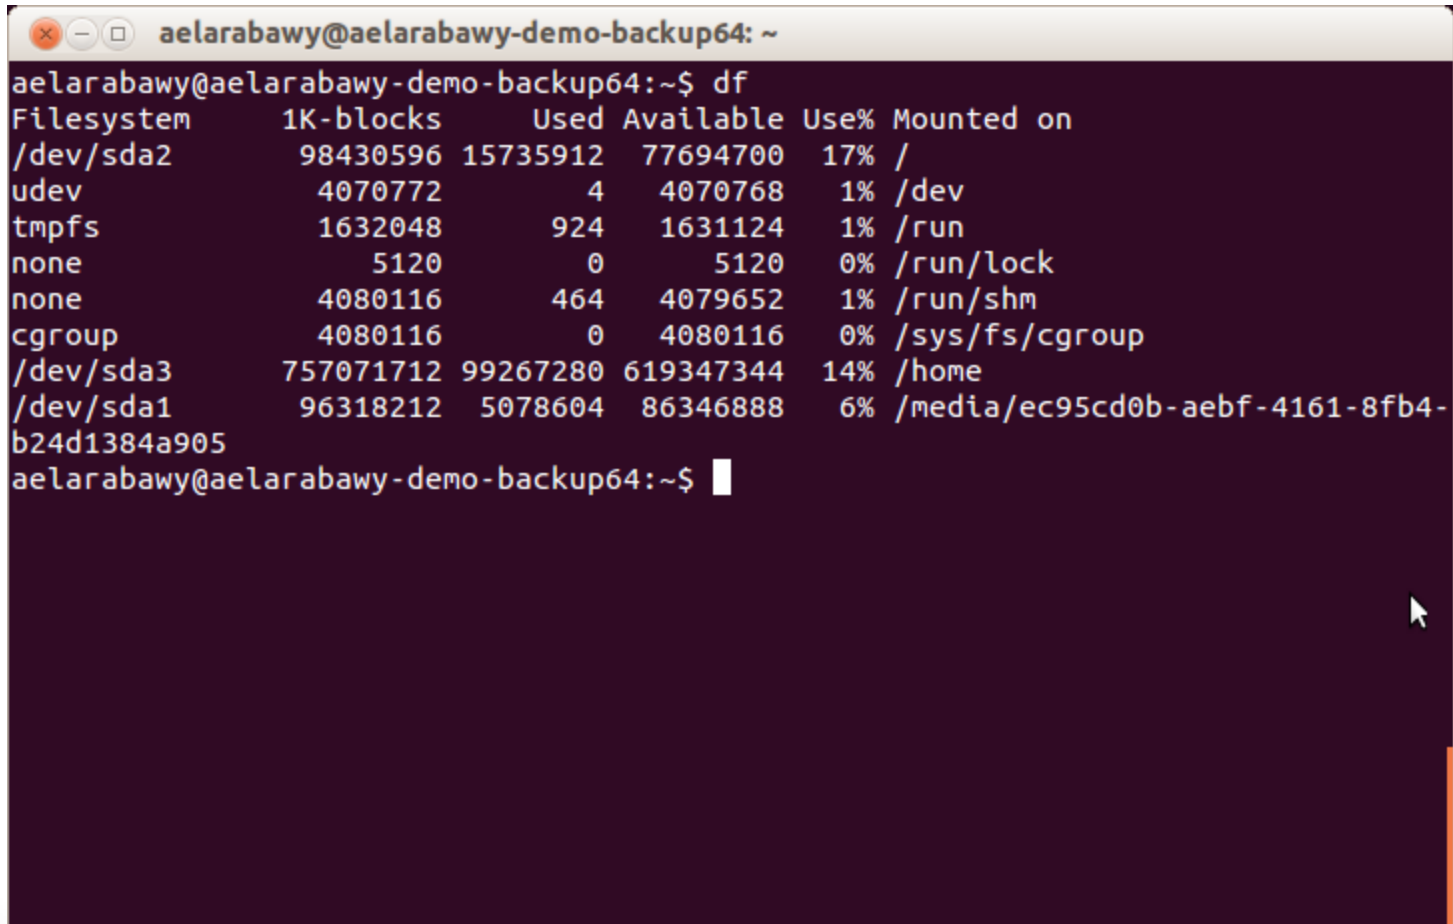
- Those filesystems are not stored in any storage device, but they are managed by the kernel
- Reading from a file in this filesystem results in a query to the kernel
- Writing to a file results in sending some info to the kernel
- Those filesystems will be studied in detail in separate lectures



LINUX COMMANDS

Show FileSystem Disk Space Usage (df Command)

\$ df



The image shows a terminal window titled "aelarabawy@aelarabawy-demo-backup64: ~". The user has entered the command "df". The output is a table showing disk space usage for various filesystems. The columns are: Filesystem, 1K-blocks, Used, Available, Use%, and Mounted on. The data is as follows:

Filesystem	1K-blocks	Used	Available	Use%	Mounted on
/dev/sda2	98430596	15735912	77694700	17%	/
udev	4070772	4	4070768	1%	/dev
tmpfs	1632048	924	1631124	1%	/run
none	5120	0	5120	0%	/run/lock
none	4080116	464	4079652	1%	/run/shm
cgroup	4080116	0	4080116	0%	/sys/fs/cgroup
/dev/sda3	757071712	99267280	619347344	14%	/home
/dev/sda1	96318212	5078604	86346888	6%	/media/ec95cd0b-aebf-4161-8fb4-b24d1384a905

The terminal prompt "aelarabawy@aelarabawy-demo-backup64:~\$" is visible at the bottom of the window.

Show FileSystem Disk Space Usage (df Command)

\$ df -i (Show FileSystem inode Usage)

```
aelarabawy@aelarabawy-demo-backup64: ~$ df -i
```

Filesystem	Inodes	IUsed	IFree	IUse%	Mounted on
/dev/sda2	6250496	363722	5886774	6%	/
udev	1017693	518	1017175	1%	/dev
tmpfs	1020029	464	1019565	1%	/run
none	1020029	5	1020024	1%	/run/lock
none	1020029	51	1019978	1%	/run/shm
cgroup	1020029	9	1020020	1%	/sys/fs/cgroup
/dev/sda3	48078848	470652	47608196	1%	/home
/dev/sda1	6119424	256011	5863413	5%	/media/ec95cd0b-aebf-4161-8fb4-b24d1384a905

```
aelarabawy@aelarabawy-demo-backup64:~$
```

Show Process Disk Usage (du Command)



\$ du <device>

- This command shows the disk usage per process for the disk specified by the device filename

\$ du /dev/sda1



Linux 4

Embedded Systems

<http://Linux4EmbeddedSystems.com>